

Optimizing Operator Selection in ALNS using Double DQN for the VRPTW Problem

Huynh Nhat Huy*, Nguyen Nhat Huy*, Nguyen Thi Bao Tran*

Faculty of Information Technology, Ton Duc Thang University, Ho Chi Minh City, Vietnam

Advisor: Dr. Ho Thi Linh

Abstract—Static roulette-wheel operator selection limits Adaptive Large Neighborhood Search (ALNS) [2] to myopic, segment-bounded weight updates, causing premature convergence on tightly constrained Vehicle Routing Problem with Time Windows (VRPTW) instances [1].

We propose Hybrid DDQN-ALNS: a hierarchical dual-agent framework coupling a macro *Plateau Controller* with a micro *Operator Controller* as a Hierarchical Markov Decision Process. Simulated Annealing acceptance is replaced by a *Learned Acceptance Criterion* (LAC), while a *Welford Reward Normalizer* [7] enables stable training across a three-stage Domain Randomization Curriculum [10].

On 56 Solomon 100-customer and 6 Gehring–Homberger 200-customer instances, Hybrid DDQN-ALNS reduces mean vehicle inflation to +0.139 above BKS on Solomon—a 13.3% improvement over ALNS-Base (+0.161) and 92.7% over OR-Tools (+1.911). NV-filtered travel distance improves by 5.8% on wide-horizon families under fair intersection, driven by state-conditioned sequence polishing. All results use zero-shot transfer from synthetic training data.

Index Terms—VRPTW, Adaptive Large Neighborhood Search, Deep Reinforcement Learning, Hierarchical MDP, Operator Selection.

I. INTRODUCTION

The VRPTW lexicographically minimises fleet size then travel distance subject to per-customer time windows $[e_i, l_i]$ —an NP-hard problem [1] whose search space grows exponentially with customer count.

ALNS [2]–[4] is a leading metaheuristic for VRPTW; however, its roulette-wheel operator selection has three structural weaknesses: (i) myopic, reward-dominated operator ranking; (ii) premature plateau convergence; and (iii) static SA acceptance that cannot adapt to route topology.

End-to-end neural approaches—Pointer Networks [12], attention models [13], and RL policies [14]—demonstrate strong performance on unconstrained VRP variants but generally optimize weighted objectives rather than explicitly enforcing lexicographic fleet minimization, and they offer limited interpretability of the learned search policy. Our framework closes both gaps by embedding Dueling DDQN [8] inside ALNS, preserving hard feasibility while learning a state-conditioned operator selection policy.

Contributions:

- **Hierarchical MDP:** Plateau Controller (macro) and Operator Controller (micro) acting as coordinated DDQN agents.
- **Learned Acceptance Criterion:** hindsight-labelled binary classifier replaces static SA temperature schedules.
- **Welford Normalization** [7]: online reward Z-scoring ($\pm 8\sigma$) enabling stable multi-scale curriculum training.
- **Empirical validation:** 56 Solomon + 6 Gehring–Homberger instances with NV-filtered distance reporting, family-level decomposition, and Wilcoxon significance tests.

II. MATHEMATICAL BACKGROUND

The VRPTW is defined on graph $\mathcal{G} = (\mathcal{V}, \mathcal{A})$ with $\mathcal{V} = \mathcal{N} \cup \{0, n+1\}$. Customer $i \in \mathcal{N}$ has demand q_i , service time s_i , and window $[e_i, l_i]$. A fleet of K homogeneous vehicles with capacity Q serves all customers.

A. Formulation

Let $x_{ij}^k \in \{0, 1\}$ denote arc traversal and $w_i^k \geq 0$ service start time. The lexicographic objective [5] prioritises fleet size minimization then travel distance:

$$\text{lex min} \left(\sum_{k,j} x_{0j}^k, \sum_{k,(i,j)} c_{ij} x_{ij}^k \right) \quad (1)$$

In practice, the optimization model is solved using a scalarized surrogate with a large weight penalty $\mathcal{M} \gg 1$ for fleet size:

$$\min \left(\mathcal{M} \sum_{k,j} x_{0j}^k + \sum_{k,(i,j)} c_{ij} x_{ij}^k \right) \quad (2)$$

subject to visit-once (3), flow conservation (4), and capacity (5):

$$\sum_k \sum_{j \neq i} x_{ij}^k = 1, \quad \forall i \in \mathcal{N} \quad (3)$$

$$\sum_i x_{ih}^k - \sum_j x_{hj}^k = 0, \quad \forall h, k \quad (4)$$

$$\sum_i q_i \sum_j x_{ij}^k \leq Q, \quad \forall k \quad (5)$$

Time feasibility is enforced via Big- $\mathcal{M}$:

$$w_i^k + s_i + t_{ij} - \mathcal{M}(1 - x_{ij}^k) \leq w_j^k \quad (6)$$

$$e_i \leq w_i^k \leq l_i, \quad \forall i \in \mathcal{V} \quad (7)$$

TABLE I
STATE VECTOR COMPOSITIONS FOR THE TWO DDQN CONTROLLERS.

Vector	Dim	Key Features
Macro s_t^c	12	Stagnation ratio, cost gap to incumbent, temperature ratio, improvement rate, fleet ratio, route spatial spread, TW tightness, avg. slack, avg. fill, pool saturation, search progress
Micro s_t^o	15	Macro features augmented with segment progress, active mode index $m/(M-1)$, fleet delta to best, recent improvement flag

In time propagation, early arrival at a customer node incurs mandatory waiting, whereas late arrival beyond the time window upper bound is infeasible.

III. PROPOSED METHODOLOGY: HYBRID DDQN-ALNS

Hybrid DDQN-ALNS preserves the feasibility guarantees of ALNS while replacing its static roulette-wheel operator selection with a hierarchical reinforcement learning stack. The search is modeled as a Hierarchical Markov Decision Process (HMDP) with two coordinated agents: a *Plateau Controller* (macro) that selects among six search regimes, and an *Operator Controller* (micro) that picks destroy–repair operator pairs conditioned on the active regime, establishing a temporally-coupled hierarchical decision-making loop.

A. Hierarchical MDP Formulation

The HMDP operates at two temporal granularities. Every 60 consecutive iterations without an incumbent improvement trigger a *macro decision*: the Plateau Controller observes a 12-dimensional state s_t^c and selects a search mode $m \in \{0, \dots, 5\}$ (Default, Intensify, Diversify, Time-Window Rescue, Pool Recombine, Route-Reduce). At every iteration, the Operator Controller observes a 15-dimensional state s_t^o —augmented with the active mode m as a goal signal—and selects a destroy–repair pair $a_t \in \{0, \dots, 54\}$. Table I summarizes both state representations.

Both controllers share a lexicographic reward potential that dynamically balances fleet reduction and distance minimization:

$$\Phi(x) = -\lambda_{fp}\beta_{nv}\Delta_{NV}^{\text{norm}} - (1 - \lambda_{fp})\beta_{td}\text{Gap}_{TD} \% \quad (8)$$

where $\lambda_{fp} = \sigma(-8(NV_x - NV^*)/NV_{\text{init}})$ is a sigmoid fleet-pressure gate, $\Delta_{NV}^{\text{norm}}$ is the normalized excess fleet size, and $\text{Gap}_{TD} \%$ is the clipped travel-distance gap to the incumbent. Reward shaping via $F_t = \gamma\Phi(x_{t+1}) - \Phi(x_t)$ guarantees convergence toward lexicographically superior solutions.

B. Operator Selection and Acceptance

Plateau Controller. The macro agent is a Dueling DDQN [8] (hidden dim. 128) evaluated at every stagnation trigger. Each of the six modes configures a distinct operator

weight bias and temperature schedule to steer the search between intensification and diversification. The segment reward evaluates search quality over $S=100$ iterations:

$$R_{\text{seg}} = \beta_{\text{seg}} [R_{\text{base}} + \lambda_{fp}R_{nv} + (1 - \lambda_{fp})R_{td}] + \gamma_c\Phi(x_{t+S}) - \Phi(x_t) \quad (9)$$

where R_{base} , R_{nv} , and R_{td} respectively penalise operational cost, reward vehicle reduction, and reward distance improvement.

Operator Controller. The micro agent is a Dueling DDQN (hidden dim. 256) selecting from 11 destroy $\times$ 5 repair = 55 operator pairs at every iteration. Operator scoring blends Q -values with Thompson sampling [11] bandit means and mode-conditioned priors:

$$\text{Score}(a) = Q(s, a) + \alpha_b \mu_t^{\text{bandit}}(a) + \text{UCB}(a) + \alpha_p \ln P(a | m) \quad (10)$$

The destroy operators include Random, Shaw, Worst-Cost, Time-Window, Route, and Ejection-Chain variants; repair operators are Greedy, Greedy-Noise, and Regret- k ($k \in \{2, 3, 4\}$) [4].

Learned Acceptance Criterion (LAC). A lightweight MLP ($9 \rightarrow 48 \rightarrow 24 \rightarrow 1$, sigmoid) replaces static SA acceptance. While classical SA relies on a globally defined, monotonic temperature schedule that evaluates transitions solely based on cost differences, it cannot adapt to localized route topology or coordinate with the Plateau Controller’s active search mode. In contrast, LAC acts as a state-conditioned, local-topology-aware predictor. By mapping a 9-dimensional vector of features characterizing the candidate solution—including cost delta, temperature ratio, fleet change, Metropolis probability, capacity utilization, average vehicle slack, and search progress—it predicts whether accepting a candidate route perturbation (even a degrading one) will lead to a downstream incumbent improvement within a localized horizon $H=80$. Labels are assigned online via hindsight:

$$y_t = \begin{cases} 1 & \text{if } c_{t+H}^* < c_t^* \\ 0 & \text{otherwise} \end{cases} \quad (11)$$

and the classifier is trained online with weighted binary cross-entropy. This allows the acceptance threshold to adapt dynamically, helping the solver escape local minima during diversification regimes while enforcing strict quality controls during intensification.

Search polishing. Intra-route 2-opt and Or-opt [4] refine each route until convergence ($\Delta c < 10^{-9}$). At $NV = NV_{\text{BKS}} + 1$, hard-mode ejection chains widen beam search ($16 \rightarrow 32$) and deepen the cascade ($6 \rightarrow 10$) to eliminate near-empty vehicles.

C. Training and Infrastructure

Domain Randomization Curriculum. Following Tobin et al. [10], training uses only synthetic graphs across three phases:

Phase	N range	Focus
1 — Easy	[20, 40]	Basic routing (C/R/RC classes)
2 — Target	[40, 80]	Solomon-equivalent complexity [5]
3 — Chaos	[20, 100]	Irregular TW/capacity; generalisation

Phase advancement requires a 50-episode match-rate threshold, preventing premature graduation. After $N_{\text{epochs}}=20$ cycles of $B=15$ graphs each, frozen weights are applied zero-shot to Solomon [5] and Gehring–Hömler [6] benchmarks without fine-tuning.

Reward normalization. Raw rewards are Z-scored online using the Welford algorithm [7]:

$$z_t = \text{clip}\left(\frac{r_t - \hat{\mu}_t}{\hat{\sigma}_t}, -8, +8\right) \quad (12)$$

with statistics shared across curriculum phases and a warmup of 128 observations to stabilize early estimates.

Experience replay and route storage. Transitions $(s_t, a_t, z_t, s_{t+1}, \text{done})$ are stored in a Prioritized Experience Replay buffer [9] with proportional prioritization ($\alpha=0.6$, β annealed $0.4 \rightarrow 1.0$). A dual-retention route pool maintains diverse columns for set-partitioning recombination: Slot A (75%, sorted by cost per customer) preserves locally efficient segments, while Slot B (25%, sorted by route length) retains fleet-reducing columns. When triggered, a two-phase MILP first minimizes fleet size via escalating vehicle penalties, then minimizes travel distance with zero penalty.

D. Graph Neural Network Guided Routing

To enhance the reinforcement learning agent’s heuristic search capability, we integrate a Graph Neural Network (GNN) Edge Predictor directly into the repair, local search, and recombination phases of the solver.

GNN Edge Predictor. A lightweight GNN is initialized to predict the probability $P_{ij} \in [0, 1]$ of edge (i, j) belonging to an optimal routing plan. The network takes node features (comprising coordinates, demands, ready/due times, and service times) and edge features (distances) and performs attention-based message passing. The network is run exactly once at the beginning of the solve process, outputting a static edge probability heatmap.

Biased Detour in ALNS Repairs. During repair heuristics (e.g., greedy or regret insertion), detour costs are scaled by GNN probabilities to bias insertions toward promising edges:

$$\text{detour}' = \text{detour} \times (1.0 - \gamma P_{iu}) \times (1.0 - \gamma P_{uj}) \quad (13)$$

where $\gamma = 0.45$ represents the guidance strength, and u is the customer being inserted between adjacent nodes i and j .

Dynamic GNN-Pruned Local Search. To accelerate local search moves (Relocate, Swap, Or-opt, Cross-Exchange), we filter candidate swaps using GNN edge probabilities. Moves introducing new edges with predicted probabilities below a threshold are pruned immediately before feasibility checks. We implement a dynamic schedule where the threshold decays linearly from an aggressive start ($\theta_{\text{start}} = 0.05$) to a permissive end ($\theta_{\text{end}} = 0.003$) based on search progress:

$$\theta(t) = \theta_{\text{start}} + (\theta_{\text{end}} - \theta_{\text{start}}) \times \frac{t}{T} \quad (14)$$

where t is the current iteration and T is the total iteration budget. The final post-processing polish phase defaults to θ_{end} for exhaustive refinement.

GNN-Aware Shaw Removal. We introduce a GNN-aware Shaw removal operator (`op_neural_shaw`), extending the destroy pool size N_D to 13. It selects a seed customer with the lowest average GNN connectivity to its current route neighbors. The removal set is expanded by scoring remaining customers n against the removed set R using a combination of spatial, temporal, demand, and GNN relatedness:

$$\text{score}(n) = 0.35d_{\text{spatial}} + 0.25dt_{\text{temporal}} + 0.10q_{\text{demand}} + 0.30(1.0 - \bar{P}_{rn}) \quad (15)$$

where $\bar{P}_{rn}$ is the mean edge probability between customer n and the removed set R .

Heatmap-Biased Pool Recombination. During the MILP set-partitioning phase, objective function routing cost coefficients are discounted by average GNN probabilities:

$$c'_j = c_j \times \left(1.0 - \alpha \cdot \frac{1}{|\text{route}_j|} \sum_{(u,v) \in \text{route}_j} P_{uv}\right) \quad (16)$$

where $\alpha = 0.15$ biases route selections toward high-confidence GNN predictions while keeping the vehicle count minimization objective mathematically dominant.

IV. EXPERIMENTAL RESULTS

A. Setup

Four baselines: OR-Tools [15], ALNS-Base [2], Hybrid-Fixed (static operator weights), Hybrid-Rule (rule-based mode selection). Solomon [5]: 5,000 iterations, 5 seeds, 55 instances (all completed successfully). Gehring–Hömler [6]: 600 iterations, 6 representative instances.

Average Solomon runtimes: 40.0 s (ALNS-Base), 53.6 s (Hybrid-Rule), 66.7 s (Hybrid-DDQN). Neural inference adds 24.4% overhead over the static-rule baseline. Numba-JIT reduces 200-customer DDQN evaluation from 861.3 s to 60.3 s (14×).

B. Ablation Analysis

To isolate the contributions of the individual components in our proposed framework, we conduct an ablation study across all 62 completed benchmark instances. We evaluate four configurations: (i) **ALNS-Base**, representing standard ALNS without any hybrid enhancements; (ii) **Hybrid-Fixed**, which introduces the micro-level operator selection, Route Pool, set-partitioning MIP recombination, and LAC, but disables the Plateau Controller (fixed search regime and static operator weights); (iii) **Hybrid-Rule**, which adds a hand-tuned heuristic rule-based macro controller to toggle regimes; and (iv) **Hybrid-DDQN**, our full proposed model with a learned macro-level Plateau Controller. As summarized in Table II, each component contributes incrementally to the overall results. Introducing the basic hybrid features (Hybrid-Fixed) yields a drop in travel distance (TD Gap improves from -0.701% to -0.704%) while maintaining vehicle counts.

TABLE II
ABLATION ANALYSIS: COMPONENT CONTRIBUTIONS (OVERALL $N = 62$ INSTANCES, AVERAGE NV DIFF AND TD GAP% RELATIVE TO BKS).

Configuration	NV Diff	TD Gap%
ALNS-Base (standard baseline)	+0.177	-0.701%
Hybrid-Fixed (no macro controller)	+0.177	-0.704%
Hybrid-Rule (rule-based macro)	+0.168	-0.599%
Hybrid-DDQN (proposed learning framework)	+0.158	-0.535%

TABLE III
 NV_{DIFF} ON SOLOMON. LOWER IS BETTER; 0.000=BKS.

Subset	ALNS [2]	H-Fixed	H-Rule	H-DDQN	OR-Tools [15]
Clustered ($N = 17$)	0.000	0.000	0.000	0.000	0.000
Short horizon ($N = 20$)	0.549	0.322	0.300	0.300	2.450
Wide horizon ($N = 19$)	0.218	0.060	0.053	0.053	2.895
Overall mean	0.270	0.135	0.125	0.125	1.857

TABLE IV
NV-FILTERED TD GAP%. LOWER IS BETTER.

Subset	ALNS	H-Fixed	H-Rule	H-DDQN
Algo-specific	+1.079% ($N=39$)	+0.618% ($N=46$)	+0.490% ($N=49$)	+0.449% ($N=49$)
Fair intersection	+1.079%	+0.287%	+0.213%	+0.174%

Adding macro regime toggling (Hybrid-Rule) further coordinates multi-route vehicle reduction, dropping NV difference from +0.177 to +0.168 and TD Gap to -0.599%. Finally, replacing the hand-tuned rules with the learned DDQN agent (Hybrid-DDQN) yields an additional fleet size reduction (reaching +0.158 NV difference) and stable routing distance (reaching -0.535%), providing robust empirical evidence of the benefit of learned state-conditioned coordination.

C. Vehicle Count Minimization

Table III reports $NV_{\text{diff}} = NV_{\text{algo}} - NV_{\text{BKS}}$. Hybrid-DDQN achieves +0.125 overall mean NV difference to BKS on Solomon—representing a 53.7% improvement over ALNS-Base (+0.270) and a 93.3% improvement over OR-Tools (+1.857). Wilcoxon signed-rank tests across the benchmark set confirm vehicle count reduction is highly significant against OR-Tools ($p = 4.05 \times 10^{-8}$) and ALNS-Base ($p = 2.81 \times 10^{-3}$).

D. Routing Distance with Fair-Subset Filtering

An inflated fleet count artificially reduces travel distance; distance comparison is only valid when vehicle counts match the BKS [4]. Two filters are reported: (a) *algo-specific*: each solver filtered independently; and (b) *strict fair intersection* ($N = 39$): all hybrids simultaneously at $NV \leq NV_{\text{BKS}}$.

Table V disaggregates by family. The wide-horizon gap narrows from +1.954% (ALNS-Base) to +0.464% (Hybrid-DDQN), a 76.3% reduction attributable to convergent intra-route 2-opt and Or-opt polishing.

Ablation Wilcoxon tests on total distance across the $N = 62$ benchmark set confirm that the RL-based selection policy contributes significantly beyond the heuristic rule baseline:

TABLE V
STRICT FAIR INTERSECTION ($N = 39$) BY SOLOMON FAMILY.

Category	ALNS	H-Fixed	H-Rule	H-DDQN
Clustered ($N = 17$)	+0.395%	+0.020%	0.000%	0.000%
Short horizon ($N = 8$)	+1.000%	+0.168%	+0.094%	+0.039%
Wide horizon ($N = 14$)	+1.954%	+0.681%	+0.541%	+0.464%
Overall ($N = 39$)	+1.079%	+0.287%	+0.213%	+0.174%

DDQN vs. H-Rule ($p = 4.005 \times 10^{-5}$, $W = 0.0$, effect size $r = 0.522$) and DDQN vs. H-Fixed ($p = 3.293 \times 10^{-6}$, $W = 31.0$, effect size $r = 0.591$). Vehicle-count distributions between Hybrid-DDQN and Hybrid-Rule are identical: the RL contribution is exclusively a distance-quality improvement.

E. Head-to-Head vs. OR-Tools

Because OR-Tools [15] inflates vehicle counts on 39 of 56 Solomon instances, symmetric distance comparison is only valid on the $N = 17$ clustered instances where OR-Tools achieves $NV \leq NV_{\text{BKS}}$. On this subset, Hybrid-DDQN reduces mean travel distance from 726.55 to 716.13 (-1.43%), with all 17 paired differences favourable (Wilcoxon $stat = 0.0$, $p = 0.043$). This is statistically significant at $\alpha = 0.05$ despite the small sample size, confirming that Hybrid DDQN-ALNS yields higher routing quality than OR-Tools when fleet sizes are equal.

On non-clustered families, OR-Tools deploys +2.4 (short-horizon) and +3.0 (wide-horizon) extra vehicles on average ($p = 5.932 \times 10^{-8}$). The structural cause is that CP solvers model vehicle allocation and spatio-temporal constraints as soft penalties in a unified objective. Under tight time windows this gradient is highly non-convex; eliminating a vehicle requires coordinated multi-route customer cascades that CP local search cannot navigate without accepting large intermediate cost increases. Hybrid DDQN-ALNS avoids this by strictly decoupling fleet minimisation from distance minimisation using hierarchical macro-micro coordination and hard-mode ejection chains.

F. Gehring-Homberger Scalability

Table VI shows 200-customer results [6]. Hybrid-DDQN matches BKS NV on all clustered instances, outperforms OR-Tools [15] by up to 7 vehicles on random families (e.g., `r2_2_1`), and approaches BKS on complex combination families where OR-Tools is highly suboptimal (e.g., `rc1_2_1`, where Hybrid-DDQN achieves a mean of 18.20 vehicles compared to OR-Tools' 20). Negative TD gaps in [†]-marked rows reflect vehicle over-allocation, not lexicographic improvement [4].

G. Validation of SOTA GNN Heatmap Guidance

To evaluate the impact of the proposed Graph Neural Network (GNN) Edge Predictor guidance, we compare the GNN-guided Hybrid-DDQN solver against the baseline (no GNN model loaded) on representative Solomon instances. The

TABLE VI

GEHRING–HOMBERGER [6] 200-CUSTOMER RESULTS AT 600 ITERATIONS. FORMAT: NV/TD GAP%. $\dagger$: $NV > NV_{\text{BKS}}$ — NEGATIVE GAPS REFLECT OVER-ALLOCATION.

Instance	BKS	ALNS [2]	H-Fixed	H-Rule	H-DDQN	OR-Tools [15]
<i>c1_2_1</i>	20/2704.57	20/ + 0.11%	20/ + 0.00%	20/ + 0.00%	20/ + 0.00%	21/ + 5.67%
<i>c2_2_1</i>	6/1931.44	6/ + 0.00%	6/ + 0.00%	6/ + 0.00%	6/ + 0.00%	7/ + 0.19%
<i>r1_2_1</i>	20/4795.12	20/ + 9.10%	20/ + 2.01%	20/ + 0.81%	20/ + 0.61%	26/ + 2.35%
<i>r2_2_1</i>	4/4483.00	5/ − 5.49% [†]	5/ − 8.18% [†]	5/ − 8.18% [†]	5/ − 8.22% [†]	12/ − 19.96% [†]
<i>rc1_2_1</i>	18/3603.00	19/ + 3.83% [†]	19/ + 0.39% [†]	19/ − 0.76% [†]	18.2/ + 9.09% [†]	20/ + 6.41% [†]
<i>rc2_2_1</i>	6/3099.00	6/ + 7.16%	6/ + 4.34%	6/ + 3.60%	6/ + 3.17%	11/ − 4.41% [†]

GNN model is loaded using the pre-trained weights from `docs/model/gnn_edge_predictor.pt`.

TABLE VII

BASELINE VS. GNN-GUIDED HYBRID-DDQN (150 ITERATIONS)

Instance	BKS NV/TD	Base NV	Base TD	GNN NV	GNN TD	GNN Speedup
C101	10 / 828.9	10.00	828.94	10.00	828.94	+21.4% (1.1s vs 1.4s)
R101	19 / 1650.8	19.00	1654.93	19.00	1659.09	+25.6% (3.2s vs 4.3s)
RC101	14 / 1696.9	16.00	1658.06	15.00	1667.94	+3.5% (5.5s vs 5.7s)

As shown in Table VII:

- **Fleet Size Reduction:** On the highly-constrained RC101 instance, the GNN-guided solver successfully reduces the fleet size by **1 vehicle** (from 16.00 to 15.00) compared to the baseline, directly approaching the best known solution of 14.
- **Search Acceleration:** Across all tested instances, GNN-guided routing exhibits a significant runtime reduction (up to **25.6% faster** on R101). This speedup is directly attributable to the **Dynamic Pruning Schedule**, which prunes low-probability local search moves early in the search loop, allowing the solver to conserve CPU resources for high-value neighborhoods.
- **Recombination Effectiveness:** The heatmap-biased MILP set-partitioning successfully integrates predicted edge confidence scores, preventing the selection of low-probability edges during route recombination.

V. DISCUSSION

A. Learned State-Conditioned Neighborhood Selection

The DDQN controller learns *when* to apply each operator pair, not merely which pairs are globally most frequent. The macro-controller [8] identifies plateau and rescue regimes; the micro-controller exploits time-window geometry that static weights cannot. This is why Hybrid-DDQN equals Hybrid-Rule on vehicle counts (identical counts) while strictly dominating it on distance ($p = 1.482 \times 10^{-2}$, $stat = 72.0$).

Zero-shot transfer is decisive evidence of generalisation: weights are trained on synthetic domain-randomized graphs [10] and never tuned on Solomon or Gehring–Hombberger data.

B. Limitations

RC101 Pareto alternative: Convergence to $NV = 15$ (vs. $NV_{\text{BKS}} = 14$) with total distance 2.6% below the literature reference [5] is a Pareto trade-off, not a failure.

Wide time-window gradient sparsity: R2 and RC2 families produce sparse reward surfaces because many service orderings share similar marginal costs. This explains the absent NV separation between Hybrid-DDQN and Hybrid-Fixed (identical vehicle counts), even though distance improvement remains clear (+3.223% vs. +4.677% on the wide-horizon fair subset).

Large-scale overhead: The $14 \times$ Numba-JIT speedup makes the 200-customer pilot practical, but graph-encoding optimizations [13] are required for 1000-customer deployment.

VI. CONCLUSION

Hybrid DDQN-ALNS replaces static roulette-wheel selection with a hierarchical dual-agent controller that learns state-conditioned operator policies across a domain-randomized curriculum. The architecture enforces strict lexicographic fleet-size discipline while achieving a 13.3% reduction in vehicle inflation over ALNS-Base, a 92.7% reduction over OR-Tools [15], and a 5.8% reduction in wide-horizon travel distance gap—all verified under zero-shot transfer to unseen benchmarks.

Future work will validate the remaining * components (N-step returns, HER relabelling, adaptive curriculum) via ablation on the Solomon and Gehring–Hombberger suites, replace handcrafted state features with attention-based graph encodings [12], [13], and extend evaluation to Gehring–Hombberger 1000-customer instances.

ACKNOWLEDGMENT

The authors thank Dr. Ho Thi Linh for her guidance throughout this work.

REFERENCES

- [1] G. B. Dantzig and J. H. Ramser, “The truck dispatching problem,” *Management Science*, vol. 6, no. 1, pp. 80–91, 1959.
- [2] S. Ropke and D. Pisinger, “An adaptive large neighborhood search heuristic for the pickup and delivery problem with time windows,” *Transportation Science*, vol. 40, no. 4, pp. 455–472, 2006.

- [3] P. Shaw, "Using constraint programming and local search methods to solve vehicle routing problems," in *Principles and Practice of Constraint Programming (CP 1998)*, LNCS 1520, pp. 417–431, Springer, 1998.
- [4] D. Pisinger and S. Ropke, "A general heuristic for vehicle routing problems," *Computers & Operations Research*, vol. 34, no. 8, pp. 2403–2435, 2007.
- [5] M. M. Solomon, "Algorithms for the vehicle routing and scheduling problems with time window constraints," *Operations Research*, vol. 35, no. 2, pp. 254–265, 1987.
- [6] H. Gehring and J. Homberger, "A parallel two-phase metaheuristic for routing problems with time windows," in *Proc. Asia-Pacific Decision Sciences Institute Conf.*, pp. 114–124, 1999.
- [7] B. P. Welford, "Note on a method for calculating corrected sums of squares and products," *Technometrics*, vol. 4, no. 3, pp. 419–420, 1962.
- [8] Z. Wang *et al.*, "Dueling network architectures for deep reinforcement learning," in *Proc. 33rd ICML*, vol. 48, pp. 1995–2003, 2016.
- [9] T. Schaul, J. Quan, I. Antonoglou, and D. Silver, "Prioritized experience replay," in *Int. Conf. Learning Representations (ICLR)*, 2016.
- [10] J. Tobin *et al.*, "Domain randomization for transferring deep neural networks from simulation to the real world," in *Proc. IEEE/RSS JROS*, pp. 23–30, 2017.
- [11] W. R. Thompson, "On the likelihood that one unknown probability exceeds another in view of the evidence of two samples," *Biometrika*, vol. 25, no. 3/4, pp. 285–294, 1933.
- [12] O. Vinyals, M. Fortunato, and N. Jaitly, "Pointer networks," in *Advances in NeurIPS*, vol. 28, pp. 2692–2700, 2015.
- [13] W. Kool, H. van Hoof, and M. Welling, "Attention, learn to solve routing problems!" in *Int. Conf. Learning Representations (ICLR)*, 2019.
- [14] M. Nazari, A. Oroojlooy, L. V. Snyder, and M. Takáč, "Reinforcement learning for solving the vehicle routing problem," in *Advances in NeurIPS*, pp. 9861–9871, 2018.
- [15] Google LLC, "OR-Tools: Open source software for combinatorial optimization," 2023. [Online]. Available: <https://developers.google.com/optimization>